

Object Detection Project: Compare STOA models to traditional algorithms in terms of mAP

DDA4220 Final Project

Xiangyi Li (119010153)

School of Data Science

Chinese University of Hong Kong, Shenzhen

xiangyili@cuhk.edu.cn

Yufei Ma (119010233)

School of Data Science

Chinese University of Hong Kong, Shenzhen

yufeima@cuhk.edu.cn

Abstract

In this project, we investigated the development of object detection algorithms in computer vision, and in the context of the VOC2012 object detection competition compare methods before and after the proposal of the R-CNN deep neural network. We also tried to reproduce STOA pre-trained models' results using PyTorch and then fine-tune the models using the VOC2012's training data. In the end we will compare the results among the traditional algorithms, the STOA algorithms, and our fine-tuned models.

1 Workload

- Xiangyi Li (50%): Responsible for implementing Faster R-CNN in PyTorch, the code framework of which will be used in other models. Responsible for setting up fine-tuning. Responsible for reading some of the paper writing.
- Yufei Ma (50%): Responsible for implementing YOLO and SSD models in PyTorch. Responsible for setting up GPUs. Responsible for writing the majority of the report.

2 Introduction

Object detection is one of the most fundamental and well-established tasks in the computer vision area, and has not been effectively solved until the re-introduction of convolutional neural networks for image classification [1]. Early works like the Viola-Jones algorithm [2], Histogram of Oriented Gradients feature descriptors [3] that mainly used handcrafted features have inspired methods like the the Hybrid Coding for Selective Search [4] that had a relatively good performance on the VOC2012 object detection competition. According to [5], after AlexNet in 2012, deep neural networks have advanced the object detection task by a large margin both in terms of accuracy and speed.

In this project, we first examined previous works introduced and approached the problem of image understanding. Then, we looked into how object detection originated from the task of face detection before deep neural network became the standard approach to dealing with computer vision tasks. Then for the VOC2012 object detection competition, we first examined how traditional object detection algorithms have performed under this task, then we went on researching and implementing STOA models and algorithms including Faster R-CNN, YOLO, and SSD. In the end, we will compare the performance among the traditional methods, the pre-trained models implemented in PyTorch or HuggingFace, and the fine-tuned models in terms of mAP, mean Accuracy Precision.

3 Related Work

Early object detection methods could date back to 2001 [2] where the Viola-Jones algorithm was first applied to face recognition, which can be considered as the predecessor to the object detection task.

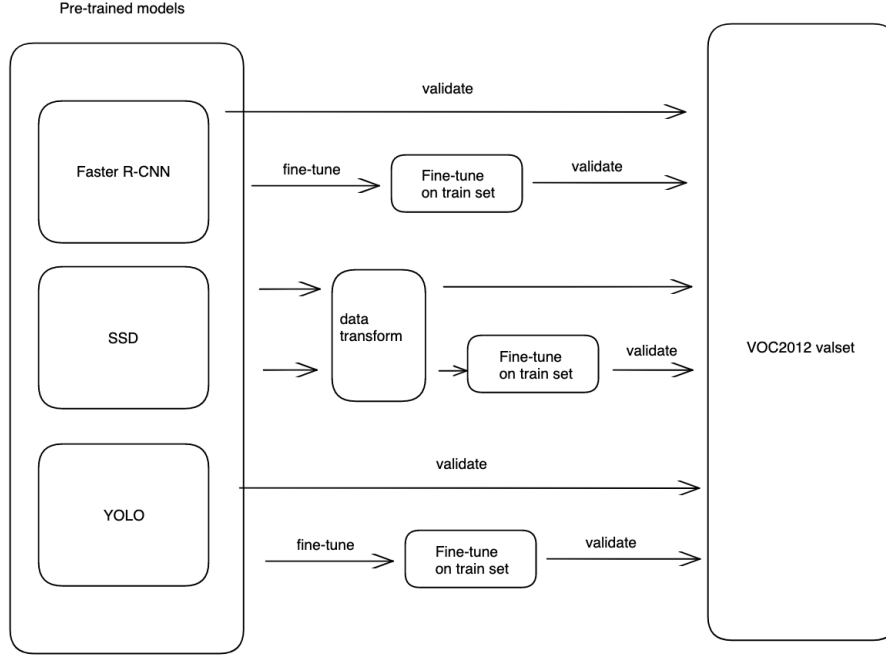


Figure 1: Designed experiment workflow for three STOA models.

It completes face recognition through a cascade of classifiers. Since its detection of side face images is unstable, different feature point representations, such as SIFT and HOG, have been explored in subsequent studies; these representations are robust to geometric and optical variations [3, 6].

Though traditional methods have seen advancements in detecting features, deep learning algorithms have gradually replaced traditional algorithms for the manual labeling of features by self-learning features. The regional CNN (R-CNN) proposed by [7], is the earliest two-stage algorithm that uses selective search methods to generate candidate regional networks and uses CNNs to extract features fed into classifiers. In 2014, SPPNet was proposed that uses different sizes of convolution kernels [8]. Fast R-CNN replaced the last layer with the pooling layer of interest [9]. Faster R-CNN proposed by [10] uses anchor boxes of different scales to extract features, reducing the generation time of candidate boxes.

Since two-stage algorithm that has high detection accuracy, numerous one-stage algorithms, such as YOLO and SSD [11, 12] have focused on making real-time object detection by boosting detection speed through techniques like elimination of region proposal step and single pass through the network. In other words, instead of first identify areas in a picture where objects might exist then look closely at each of the identified areas in the two-stage algorithms, one-stage algorithms, for example YOLO, splits the pictures into a grid where each cell in the grid predicts if it contains an object and what that object is. This is done in a single pass through the network, which makes it much faster.

According to a survey published in 2019 [13], other challenges such as robustness to scales, orientations, and occlusion and clutter are still met in the object detection task. Common techniques seen to deal with this include data augmentation with generative adversarial networks, RNN, and Transformers have been introduced to facilitate the dealing with occlusion and clutter [14, 15].

4 Method

For the three models we explored and investigated, we have implemented transforming, training loop, evaluation metric, and model fine-tuning from scratch in PyTorch as well as used Computer Vision toolkits like OpenMMLab and GluonCV. We also used standalone implementations for different models by forking their code to adjust our needs. We first implemented the Faster R-CNN from scratch in PyTorch, got score of the pre-trained model’s performance on the VOC2012 validation

image set, and then got the score of the fine-tuned model trained on the train image set. We referenced [7] to use results prior to R-CNN as our baseline, that is mAP less than 33.3%. Also according to [12], in our attempt to use GluonCV to implement SSD we will also have to perform data transform for data augmentation in before training.

5 Experiments

5.1 Evaluation method

We used mAP, mean average precision in evaluating the performance of models in the object detection task. It is defined as the average of the precision values at different recall levels. Precision is the ratio of true positives to all predicted positives, and recall is the ratio of true positives to all ground truth positives. According to [5], mAP was first introduced in the object detection task in 2007 by PASCAL VOC challenge.

5.2 Experimental details

In planning finishing this project, we intended to go over the history of the object detection task, and tried to implement from the first object detection algorithm to the most STOA models. It turned out the work before R-CNN was very hard to replicate for two students with little computer vision knowledge, so instead we went straight to deep learning methods, which eliminated the trouble of manual feature engineering for us.

In the end we successfully complete the VOC2012 competition with three STOA models [10, 12, 16]. with different open-source pre-trained weights in different manners.

Faster R-CNN: We used PyTorch's built-in faster R-CNN pre-trained model and fine-tuned it on the training image set. The code file is `frcnn_first_try.ipynb`.

SSD with GluonCV: We explored the GluonCV toolkit and attempted to first run the validation on its built-in SSD model, but the val loop was taking too long. We surmised it was due to incompatible GPU or torch visions on instances on both the Colab and Vast.ai, since it does not support PyTorch 2.0 and we had to downgrade a lot of dependencies while encountering some cryptic error messages when doing it. Setting the batch size to be more than 1 also leads to failure. In the end, we chose not to pursue this approach. The code file is `SSD.ipynb`.

SSD with ssd.pytorch: This is an re-implementation of the original SSD model in 2016 developed by a group of developers as a side project. The implementation used a very early version of PyTorch, around 1.3.0, of which a lot of functionalities and APIs have been deprecated, with many unresolved GitHub issues that led us to troubles when try to run the model on Colab. Fortunately there was abundant discussion in the community, and we were able to fork the repository to incrementally modify the code so that it became compatible with PyTorch 2.0 in the end, which is also the version of PyTorch that Colab runs. For this approach we tested their pre-trained model on the VOC2012 and got a decent score of mAP. Then we also tried to fine-tune a pre-trained model based on the COCO dataset for three epochs that got a not so good result.

YOLOv3 with GluonCV: similar scenario to SSD with GluonCV. YOLOv5 with Ultralytics: This was the simplest task since we've already figured out all the setbacks (see below) in this project, and were able to smoothly train the YOLO model by setting up scripts on GitHub. For this model, we also got the most detailed analytics that are generated by scripts during the training process.

In the process we also have had numerous failures due to various reasons. These are listed as below:

Google Colab: This was our go-to choice since we completed assignments on it and it could access advanced computing power like Nvidia A100 GPUs. However, the relationship between Colab's virtual machines' disk spaces and Google Drive is very complicated, which eluded our knowledge before we spent countless hours on it. We were misguided by many previous tutorials which often mount their execution space to Google Drive at the beginning of their file. This led to very encrypted bugs when performing intensive I/O operations due to the nature of this task where constant read & write from the annotation files and image files were necessary. It took us a few more setbacks on other cloud providers to realize that the error lied with using Google Drive.

Vast.ai: We rented bare-metal like machines on this platform after seeing Colab produces too many error messages during training and validating loops. It turned out to be a worse choice for a few reasons. First is that the VM instances are not as standardized as the Colab VM, it led us to disk shortage, incompatible CUDA version with CV tools we were trying to install. Second is the bandwidth of the virtual server is not stable, which caused significantly longer downloading datasets and pre-trained weights time. Finally is that the nohup command does seem to work on the VM due to unknown reasons. This led us to seek a few other cloud service providers before turning back to Colab.

Vultr.com: The problem lied at an even earlier stage. We managed to deposit around \$70 in total to deploy a GPU compute instance but failed to. We contacted the support team and no solutions have been working until the time of this writing.

AWS, GCP, Azure: We encountered authentication with the storage buckets when instantiating EC2-g4dn instances on AWS. For GCP and Azure, we encountered problems with billing and could not successfully boot up any virtual machines.

Since in assignments we were exposed to the OpenMMLab [17], which is a computer vision toolkit framework like GlueCV [18] that is integrated with datasets, evaluation methods, and pre-trained models, we tried to re-implement all three algorithms using MMDetection [17] so that we could enhance our understanding of the assignments. However, we encountered problems building mmdet on both the Colab and Vast.ai. But after the experiments, we've found directly reproducing standalone STOA models' implementations is much easier and more efficient than implementing within a CV toolkit framework.

5.3 Results and analysis

Table 1: Object detection performance comparison

Method	mAP	Remark
Faster R-CNN	0.7412	PyTorch built-in pretrained model fasterrcnn_resnet50_fpn_finetuned.pth
Faster R-CNN fine-tuned	0.7253	Based on the above baseline model
SSD pretrained	0.7749	github.com/amdegroot/ssd.pytorch
SSD	N/A	Pre-trained model from GlueCV ssd_512_resnet50_v1_voc
YOLOv5	0.7010	Ultralytics official YOLOv5 implementation at github.com/ultralytics/yolov5
YOLOv5 with fine-tuning	0.7970	Ultralytics official YOLOv5 implementation at github.com/ultralytics/yolov5
YOLOv3 from GlueCV	N/A	Pre-trained model from GlueCV yolo3_darknet53_voc
MMDetection	N/A	Tried to use this to implement all three models. However, encountered numerous issues with dependencies and hardware on many cloud providers.

The result is shown in Table 1. Here's a brief summary on the results. In our experiment, a higher mAP indicates a better performance of the model.

Faster R-CNN: mAP of 0.7412 with PyTorch built-in pretrained model and 0.7253 when fine-tuned on VOC2012 train image set.

SSD pretrained: mAP of 0.7749 using github.com/amdegroot/ssd.pytorch. We didn't get results for the pre-trained model from GlueCV.

YOLOv5: mAP of 0.7010 with Ultralytics official YOLOv5 implementation and 0.7970 with fine-tuning. We didn't get results for YOLOv3 from GlueCV.

Besides that all of the models we explored achieved better results than the origin. We found out the following from the test results:

Faster R-CNN: The PyTorch built-in pretrained model achieved an mAP of 0.7412. However, after we fine-tune the model on the VOC2012 training set, the mAP slightly decreased to 0.7253. This could suggest that the fine-tuning process might not have been optimized properly, for example the number of epochs might not be enough, or the pre-trained model has already achieved the ceiling of Faster R-CNN on this dataset. It's also possible that overfitting occurred during fine-tuning, leading to a decrease in performance on the validation set.

SSD pretrained: The SSD model, when pretrained and implemented using the code from github.com/amdegroot/ssd.pytorch, achieved an mAP of 0.7749, which is higher than both the pretrained and fine-tuned Faster R-CNN models. This suggests that the SSD model might be more effective for this particular task. Unfortunately, we were unable to get results for the pre-trained model from GluonCV. We also tried to train the SSD from scratch based on a COCO pre-trained model, but it might be due to that we only trained for 3 epochs, the mAP is only 0.2023.

YOLOv5: The Ultralytics official YOLOv5 implementation achieved an mAP of 0.7010. However, when fine-tuned, the mAP significantly increased to 0.7970. This is the highest mAP among all the models and configurations tested, suggesting that the fine-tuned YOLOv5 model was the most effective for this task. We didn't get results for YOLOv3 from GluonCV, so a comparison couldn't be made. Also note that we only fine-tuned the YOLOv5 for 3 epochs to achieve such performance. By the fact that the COCO pre-trained SSD model only got 0.2023 on mAP after three epochs, it is obvious that selecting an appropriate pre-trained model is essential for object detection in terms of both accuracy and training efficiency.

6 Conclusion

This project compared the performance of several implementation of STOA methods in the object detection task, including Faster R-CNN, SSD, and YOLOv5, with the mAP being the evaluation metric. The fine-tuned YOLOv5 model achieved the highest mean Average Precision (mAP) of 0.7970, outperforming both the Faster R-CNN and SSD models. This suggests that YOLOv5, when properly fine-tuned, can be highly effective for object detection tasks. The main difficulties was dealing with the I/O system on Google Colab's virtual machine, namely unawareness the unwieldiness of using the Google Drive as the file system to perform object detection tasks.

In verifying that STOA methods in the object detection task outperform traditional methods and early detection models, we have learned how to evaluate the performance of a detection model, how to construct training and validating loops in PyTorch for the fine-tuning of pre-trained models, how to use CV toolkits like GluonCV to construct pipelines that could provide boilerplate codes and APIs that could simplify the workflow, and also how to run and modify standalone implementations of STOA models in PyTorch to train and predict. As a side shot, we also learned a great deal about the cloud computing platforms especially Colab, cloud file systems, and how to utilize the CUDA GPUs correctly and effectively (our initial training and validating of the Faster R-CNN model failed to utilize the GPU because we didn't use the correct PyTorch API, the slowdown was immense). I think the learning of the tools we are using is as important as learning object detection itself since they are equally transferable if we are to be working with other deep learning projects in the future.

There are many limitations to this project. First, we faced a significant constraint in terms of computational resources. We were unable to fully utilize GluonCV and MMDetection due to their strict dependency settings, which required more advanced computing devices than we had available. This limitation restricted our ability to experiment with these powerful tools and potentially impacted the performance of our models. Second, we did not incorporate some of the newer models such as YOLOv7 into our project. This was due to we was progressing from the earliest models and after encountering setbacks from YOLOv3 using the GluonCV toolkit, we were unsure about the stability and reproducibility of even newer models. Third, our project did not fully capture the breadth of novel ideas and approaches in the field of computer vision, which is a highly active area of research. Many new methods and improvements have emerged since the publication of all three models we examines, such as Faster Faster R-CNN for Faster R-CNN, Improved SSD Network for SSD, and YOLOs (You Only Look One Sequence) for YOLO. Mathematically, the impact of different batch size and optimization algorithm could be huge, which we didn't investigate as well.

References

- [1] Syed Sahil Abbas Zaidi, Mohammad Samar Ansari, Asra Aslam, Nadia Kanwal, Mamoon Asghar, and Brian Lee. A survey of modern deep learning based object detection models, 2021.
- [2] P. Viola and M. Jones. Rapid object detection using a boosted cascade of simple features. In *Proceedings of the 2001 IEEE Computer Society Conference on Computer Vision and Pattern Recognition. CVPR 2001*, volume 1, pages I–I, 2001.
- [3] N. Dalal and B. Triggs. Histograms of oriented gradients for human detection. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, volume 1, pages 886–893 vol. 1, 2005.
- [4] Koen van de Sande, Jasper Uijlings, Cees Snoek, and Arnold Smeulders. Hybrid coding for selective search. In *PASCAL VOC Workshop*. University of Amsterdam, 2012.
- [5] Zhengxia Zou, Keyan Chen, Zhenwei Shi, Yuhong Guo, and Jieping Ye. Object detection in 20 years: A survey, 2023.
- [6] David G. Lowe. Distinctive image features from scale-invariant keypoints. *Int. J. Comput. Vision*, 60(2):91–110, November 2004.
- [7] Ross Girshick, Jeff Donahue, Trevor Darrell, and Jitendra Malik. Rich feature hierarchies for accurate object detection and semantic segmentation, 2014.
- [8] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Spatial pyramid pooling in deep convolutional networks for visual recognition. In *Computer Vision – ECCV 2014*, pages 346–361. Springer International Publishing, 2014.
- [9] Ross Girshick. Fast r-cnn, 2015.
- [10] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks, 2016.
- [11] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection, 2016.
- [12] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C. Berg. SSD: Single shot MultiBox detector. In *Computer Vision – ECCV 2016*, pages 21–37. Springer International Publishing, 2016.
- [13] Li Liu, Wanli Ouyang, Xiaogang Wang, Paul Fieguth, Jie Chen, Xinwang Liu, and Matti Pietikäinen. Deep learning for generic object detection: A survey, 2019.
- [14] Zitong Yu, Yuming Shen, Jingang Shi, Hengshuang Zhao, Yawen Cui, Jiehua Zhang, Philip Torr, and Guoying Zhao. Physformer++: Facial video-based physiological measurement with slowfast temporal difference transformer, 2023.
- [15] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A. Efros. Image-to-image translation with conditional adversarial networks, 2018.
- [16] Joseph Redmon and Ali Farhadi. Yolov3: An incremental improvement, 2018.
- [17] Kai Chen, Jiaqi Wang, Jiangmiao Pang, Yuhang Cao, Yu Xiong, Xiaoxiao Li, Shuyang Sun, Wansen Feng, Ziwei Liu, Jiarui Xu, Zheng Zhang, Dazhi Cheng, Chenchen Zhu, Tianheng Cheng, Qijie Zhao, Buyu Li, Xin Lu, Rui Zhu, Yue Wu, Jifeng Dai, Jingdong Wang, Jianping Shi, Wanli Ouyang, Chen Change Loy, and Dahua Lin. Mmdetection: Open mmlab detection toolbox and benchmark, 2019.
- [18] Jian Guo, He He, Tong He, Leonard Lausen, Mu Li, Haibin Lin, Xingjian Shi, Chenguang Wang, Junyuan Xie, Sheng Zha, Aston Zhang, Hang Zhang, Zhi Zhang, Zhongyue Zhang, Shuai Zheng, and Yi Zhu. Gluoncv and gluonnlp: Deep learning in computer vision and natural language processing, 2020.